
Speaking Jewels into Spacetime: A Concept-and-Feasibility Study of Native Text-to-Video Gaussian Programs

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Text-to-video systems usually generate a raster latent and decode it into frames. We
2 investigate a different target language: a video is cast directly as an irregular field of
3 persistent, anisotropic Gaussian Jewels in displayed spacetime (u, v, t) . Each Jewel
4 carries geometry, opacity, color, and a local color derivative; time-tilted support rep-
5 represents motion. The intended generator speaks a hierarchical executable program—
6 semantic scene and ownership tokens, trajectory phrases, physical Jewel tokens,
7 and continuous centroids—that an additive renderer turns into video. This paper
8 makes a bounded feasibility claim, not an open-vocabulary generation claim. In
9 controlled experiments, free spacetime tilt improves reconstruction in 9/9 matched
10 runs by 0.772 dB (95% CI [0.573, 0.971]); a support-correct amortized encoder
11 improves from 21.876 to 22.628 dB as training data grows from 12 to 120 videos;
12 and a prompt-and-seed-only exact compiler generates nine 72k-Jewel programs
13 with correct prompt similarity above shuffled and null controls in 9/9 cases and
14 OpenCLIP top-1 in 8/9. A 0.54M-parameter autoregressive speaker learns the
15 program syntax and recognizes all three held-out paraphrase classes, but fails the
16 stricter rendered retrieval gate. The remaining bottleneck is explicit: current trajec-
17 tory macros are source-backed. The next conclusive experiment is to learn reusable,
18 source-independent trajectory prototypes at larger data and compute scale.

19 1 Introduction

20 A video can be treated as a stack of pixels, or as a continuous volume whose axes are horizontal
21 position, vertical position, and time. We ask whether a text-conditioned model can generate the
22 second object directly. Our primitive is a spacetime Gaussian Jewel: a colored, soft-support basis
23 function living in the displayed-video volume (u, v, t) . A slanted Jewel intersects consecutive time
24 slices at different image positions, so motion is part of its geometry rather than a correspondence
25 computed after frames are generated.

26 This target is neither a codec nor a claim of compactness. Video Gaussian representations have
27 already been studied for reconstruction, processing, and compression [16, 15, 18, 6, 14]. Our question
28 is different: can text select and compose persistent Gaussian operations as the *native output* of a
29 generative model? Such a model would speak into a continuous volume, keep explicit state across
30 temporal windows, and eventually permit local edits by modifying a bounded set of primitives.

31 The hypothesis is attractive but easy to overstate. Reconstructing an existing video proves that the
32 representation can hold video, not that language can generate it. A prompt-conditioned scaffold fol-
33 lowed by an encoder proves semantic preservation, not native prompt-to-Jewel generation. Retrieval




Figure 1: The proposed native text-to-video language. Solid boxes are demonstrated in this work; the dashed box is not. The exact compiler and learned speaker select a semantic plan, persistent foreground/background ownership, addressed phrases, and physical tokens. The crucial limitation is visible in the architecture: current macro tokens are backed by 18 fitted sources.

of a whole fitted source field is also not generation. We therefore use input-ownership audits and shuffled/null controls to isolate a narrower feasibility statement.

Figure 1 shows both the demonstrated path and its missing component. At inference, the demonstrated prompt-only path receives exactly a prompt and integer seed, produces a typed plan, composes distinct foreground and background trajectory donors, expands local physical tokens at continuous centroids, and renders an irregular 72k-Jewel field. However, its macro dictionary is built from 18 fitted training fields. It does not retrieve a complete video, but it does reuse source-backed trajectory programs. Replacing those macros with reusable, source-independent prototypes is the central next gate.

Our contributions are:

- a 22-parameter additive spacetime Gaussian primitive and a support-complete renderer that preserve continuous, irregular, time-tilted structure;
- causal and scaling evidence that temporal tilt is useful and that amortized video-to-Jewel prediction improves with data rather than collapsing to a coordinate grid;
- a hierarchical physical vocabulary with exact continuous centroids, three factorized $K = 1024$ token books, and measurable same-source canonicity;
- a bounded prompt-and-seed-only test of this pipeline, including exact and learned speakers, complete input audits, negative controls, and an explicit failure of the learned speaker’s strict rendered retrieval gate.

2 Related work

Dynamic and video Gaussians. 3D Gaussian Splatting represents radiance fields with explicit anisotropic primitives [5]. Dynamic extensions model time by deformation or native spacetime support [20, 7]. Video-oriented methods use Gaussians for processing, representation, or compression [16, 15, 18, 6, 14]. These works establish that Gaussians can represent changing visual signals. We focus on a displayed 2D video volume, additive rendering, and language-driven construction rather than novel-view synthesis or rate-distortion performance.

Generative structured media. Video diffusion generates dense latent tensors [3, 1]. Autoregressive multimodal models show that a language-modeling interface can emit visual tokens [2, 19], while Transfusion combines discrete language modeling with continuous visual generation [23]. Gaussian generation has been explored for 3D and dynamic scenes [21, 13, 8, 11]. Our proposal is complementary: the output tokens expand into persistent local operations in displayed spacetime. The contribution is not next-token prediction for media, but an empirical test of whether a hierarchical, executable Gaussian language has the required representational and semantic properties.




Figure 2: One spacetime Gaussian Jewel. A video frame is a fixed- t slice. The cross-section of tilted support changes position across slices, representing motion continuously. The color Jacobian permits local appearance change through space and time.

3 Method

3.1 A Jewel is a local function in displayed spacetime

Let the normalized video domain be $\Omega = [-1, 1]^3$, with query $\mathbf{x} = (u, v, t) \in \Omega$. Jewel i has center $\mu_i \in \mathbb{R}^3$, symmetric positive definite covariance $\Sigma_i \in \mathbb{R}^{3 \times 3}$, opacity logit ℓ_i , base RGB $\mathbf{c}_i \in \mathbb{R}^3$, and local color Jacobian $\mathbf{J}_i \in \mathbb{R}^{3 \times 3}$. We evaluate

$$q_i(\mathbf{x}) = (\mathbf{x} - \mu_i)^\top \Sigma_i^{-1} (\mathbf{x} - \mu_i), \quad (1)$$

$$a_i(\mathbf{x}) = \sigma(\ell_i) \exp[-q_i(\mathbf{x})/2], \quad (2)$$

$$\hat{\mathbf{I}}(\mathbf{x}) = \mathbf{b} + \sum_i a_i(\mathbf{x}) [\mathbf{c}_i + \mathbf{J}_i (\mathbf{x} - \mu_i)], \quad (3)$$

where $\mathbf{b}$ is a learned constant background. This is additive accumulation, not ordered alpha compositing. The primitive has 22 continuous parameters: 3 center, 6 covariance, 1 opacity, 3 base-color, and 9 first-order color parameters. Values outside the display range are clipped only for visualization; range violations are audited separately.

Writing $\Sigma_i = \mathbf{R}_i \text{diag}(s_i^2) \mathbf{R}_i^\top$ makes motion geometric. Off-diagonal space-time covariance tilts support into t ; the intersection with successive fixed- t planes moves across (u, v) . Figure 2 illustrates this without the misleading picture of a flat frame translating laterally through image space.

3.2 Support-correct irregular rendering

The original fast renderer selected a fixed number of Euclidean nearest centers. That shortcut can silently omit wide or rotated Gaussians and confound a representation test with an approximation artifact. Our support-complete renderer tiles query voxels, finds every primitive whose exact Mahalanobis support lies within five standard deviations, and raises an error rather than dropping contributors when capacity is exceeded. At 10k, 45k, and 72k primitives it runs within roughly $2\times$ the k-nearest renderer; at 72k its measured peak allocation is 8.01 GB on an RTX 4090. Pixels and gradients agree with dense evaluation up to the declared support truncation.

Irregularity is equally important. Centroids are continuous and jittered; local phrases do not place one primitive on each fixed cell. We report the fraction of centers locked to the implied grid and require zero. This separates a useful structured address from visible coordinate quantization.

3.3 A hierarchical executable language

A flat sequence of 72k independent 22-D records is too long and has no semantic ownership. We use four levels:

1. a *scene token* chooses the semantic family;
2. foreground and background *owner tokens* keep coherent donors along a full space–time trajectory instead of switching independently per cell;
3. addressed local phrases specify occupancy and which owner supplies each region; and
4. three physical token books independently quantize covariance, surface/base color, and color gradient, while centers remain continuous.

The factorization has 3×1024 active physical entries rather than a single combinatorial codebook. An eight-frame field requires 35,265 discrete decisions, below a preregistered 40k-token budget, and expands deterministically into an irregular field. The active representation reaches 22.866 dB token-only reconstruction PSNR, retains $1.042\times$ the teacher’s median mixed space–time tilt, and has zero grid locking. Same-source physical token overlap is 0.452 versus 0.228 for different sources, a 0.224 canonicity margin.

Persistent ownership emerged as the load-bearing abstraction. A global scene decoder and a purely addressed $16 \times 16 \times 8$ local phrase model each produced zero recognizable classes, although the local model improved token NLL from 6.123 to 5.168. Holding one coherent source through the volume recovered 3/3 classes. Assigning one foreground donor and a different background donor along a smoothed saliency trajectory preserved 3/3 and created a composition absent from any single source. Thus locality improves syntax, but trajectory-level ownership carries semantics.

3.4 Exact and learned speakers

The *exact compiler* maps three exact prompts and a seed to (scene, foreground owner, background owner, path). For each prompt, it chooses two distinct donors from six fitted fields in that class, with exactly 50/50 primitive ownership. The source choice is made once per temporal window, not independently per voxel. No target video, target field, block program, or held-out latent is provided at inference.

The *learned speaker* tests whether the plan has learnable syntax rather than only hand-written dispatch. A 541,223-parameter autoregressive network conditions on a 512-D text embedding and emits scene, foreground, then background tokens. It trains on 216 programs spanning each exact prompt and two paraphrases per class; an unseen fourth phrasing is reserved for evaluation. This is a small speaker over a closed three-scene vocabulary, not an LLM and not an open-vocabulary model.

4 Experimental design

We organize experiments as falsifiable gates rather than a cumulative demo. Every positive semantic claim includes a cyclic-shuffled prompt and a null-text control. Conditions share seeds and, where applicable, weights and targets. Prompt similarity is measured with OpenCLIP ViT-B/32 [12]; reconstruction uses PSNR and LPIPS [22]. CLIP is used as a screening metric, not as evidence of human-perceived video quality.

Causal representation test. We fit free full-covariance Jewels and a matched control whose space–time covariance is projected to zero after every update. Primitive count, parameter bytes, initialization family, and optimization budget are matched across three source videos and three seeds. We report paired differences and a two-sided Student-*t* 95% interval over the nine pairs.

Data scaling test. An amortized 5.26M-parameter encoder is trained on exact nested subsets of 12, 60, or 120 examples from five styles and 12 actions. Each budget has three independent seeds and is evaluated on the same frozen 60-video inventory. Models receive 120 matched corpus passes and a common low-rate continuation. Confidence intervals are *t* intervals over the three seed means; no arm is selected at a different exposure count. The fitted per-video teacher is an empirical ceiling, not a comparable generalizing model.

Prompt-only test. The exact compiler is evaluated on ballerina, dog, and welder prompts, three seeds each. Each output contains 72k Jewels and renders 49 frames at 144×216 . Success requires that correct-prompt similarity exceed both shuffled and null controls for every program, with at least

Table 1: Evidence ledger. Pass means the declared bounded gate passed; it does not extend the scope stated in the final column. Exact per-run values are retained in machine-readable reports.

Question	Control	Result	Gate	What it establishes
Does time tilt help?	Full covariance vs. zeroed space-time terms; matched resources	Free wins 9/9; +0.772 dB, 95% CI [0.573, 0.971]	pass	Space-time covariance is causally useful on the tested reconstructions.
Does the encoder scale?	Nested 12/60/120 sets, 3 seeds, fixed 60-video validation	21.876 → 22.399 → 22.628 dB; LPIPS 0.473 → 0.465 → 0.392	pass	More data improves held-out prediction under matched protocol.
Is the physical language irregular and faithful?	Token-only decode; grid-lock and tilt audits	22.866 dB; zero grid lock; 1.042× tilt; 0.224 canonicity margin	pass	Factorized tokens preserve physical structure with exact centroids.
Can prompt + seed select video?	9 correct/shuffled/null programs; no target artifact at inference	Correct beats both controls 9/9; strict top-1 8/9; all class majorities	pass	A bounded source-backed native program is prompt-selective.
Can a small speaker generalize phrasing?	Unseen fourth paraphrase; frozen checkpoint; same render audit	Scene accuracy 3/3; NLL 2.201 vs. 4.387 shuffled; rendered top-1 4/9	partial	Program syntax is learnable; rendered open-ended generalization is not established.

140 6/9 strict OpenCLIP top-1 outcomes and majority success for all classes. The learned speaker uses
 141 the same render audit on unseen paraphrases; its weights are frozen at the best held-out correct-text
 142 NLL checkpoint before generating videos.

143 5 Results

144 5.1 Temporal tilt is useful, not decorative

145 The free model beats the projected control in all nine paired fits. The mean paired PSNR gain is
 146 0.77197 dB (sample standard deviation 0.25858 dB; 95% CI [0.57321, 0.97072]). Median mixed
 147 space-time tilt is 0.3459 in the free models and exactly zero by construction in controls. Gains
 148 are larger on higher-motion sources. Because primitive counts and serialized parameter bytes are
 149 matched, the result is evidence for the orientation of support, not extra capacity. It does not prove
 150 that every video benefits or that the current parameterization is optimal.

151 5.2 The amortized bottleneck improves with data

152 Figure 3 reports the predeclared scaling study. Mean held-out PSNR improves at both budget
 153 increases: 21.876 dB (95% CI [21.751, 22.000]), 22.399 dB ([22.344, 22.453]), and 22.628 dB
 154 ([22.459, 22.797]). Every 120-example seed beats every 60-example seed. From 60 to 120 examples,
 155 LPIPS falls 15.8%, layout PSNR rises from 30.726 to 31.533 dB, and median mixed tilt grows 2.58×
 156 from 0.0725 to 0.1875. The fitted teacher ceiling is 27.699 dB, LPIPS 0.1720, and median tilt 0.3966,
 157 leaving substantial headroom.

158 The same encoder preserves prompt information through a conservative text-to-pretrained-video-
 159 to-encoder route: correct prompt similarity is 0.3339 versus 0.1427 shuffled and 0.2096 null on
 160 12 held-out actions. Correct beats shuffled in 12/12 cases and retains 91.4% of the source video’s
 161 prompt alignment. This is useful evidence that the representation does not erase semantics, but it is
 162 deliberately excluded from the native-generation claim because a pretrained video model supplies the
 163 source video.

164 5.3 The exact prompt-only compiler passes its bounded gate

165 Across nine programs, mean OpenCLIP similarity is 0.18173 for the correct prompt, 0.12525 for the
 166 cyclic-shuffled prompt, and 0.14825 for null text. The paired margins are +0.05648 and +0.03348;




Figure 3: Support-correct encoder scaling. Left: held-out PSNR against training-set size, mean and 95% t interval over three seeds. Right: convergence against matched corpus exposure. The trend supports a data-and-compute scaling hypothesis but remains a small-data study.




Figure 4: Prompt-only proof sheet reflowed from one audited image. Each half shows correct-prompt and cyclic-shuffled generation for ballerina, dog, and welder classes. All evaluated samples are shown. These are 72k-Jewel renders, not source frames; their trajectory macro tokens are nevertheless backed by fitted training fields.

167 correct exceeds each control for all 9/9 programs. Strict top-1 is 8/9 and majority retrieval passes
 168 for all three classes. Foreground and background donor identifiers are distinct in every program,
 169 ownership is exactly balanced, and grid-lock fraction is zero. The input audit confirms that the only
 170 semantic inputs are exact prompt text and the declared integer seed.

171 Figure 4 shows all conditions rather than selected successes. The exact compiler is visibly class-
 172 selective and composes donor regions, but quality is far below a modern dense video generator.
 173 Because the macro vocabulary is source-backed, this result establishes execution and semantic
 174 selection, not novel open-vocabulary motion.

5.4 The learned speaker exposes the next bottleneck

The learned speaker reaches its best held-out correct-text NLL at step 100 and continues to step 1100 before the declared plateau rule stops training. On unseen paraphrases, macro token NLL is 2.2005 for correct text, 4.3870 for shuffled text, and 2.5391 for null text. Scene accuracy is 100% for correct text, 0% shuffled, and 33% null; all nine correct-text samples are scene-consistent. In rendered space, correct similarity exceeds shuffled by 0.04782 on average and null by 0.04117, with 7/9 correct-over-shuffled pairs.

The strict rendered top-1 test nevertheless passes only 4/9, below the declared 6/9 threshold, and only one of three class majorities passes. This failure is informative but not an impossibility result. The network learned the coarse scene syntax under unseen wording; it did not reliably choose fine donor combinations whose low-fidelity renders win a frozen semantic retrieval test. More steps alone are not the leading intervention because held-out correct NLL plateaued early. The next experiment should change the vocabulary and supervision: learn reusable trajectory prototypes from many sources, train the speaker against rendered semantic and temporal objectives, and compare against a parameter-matched direct continuous head.

6 What should scale next?

The experiments isolate three different scaling questions.

Representation compute. Longer local appearance optimization materially improves the frozen-geometry derivative head. For one seed, LPIPS falls from 0.717 at 400 updates to 0.670 at 4k and 0.659 at 12k; a second seed reaches 0.660 at 12k. The final seeds differ by only 0.00058 LPIPS and 0.00041 dB PSNR, and geometry remains bit-identical. This supports additional compute for appearance refinement, but diminishing returns after roughly 8–12k updates argue against presenting optimization duration as the main missing idea.

Encoder data. The monotone held-out curve and remaining teacher gap support a larger paired video/Jewel corpus. The current 120-example maximum is far too small for a foundation model claim. A useful next scale study should fix architecture and corpus passes, increase data by at least another order of magnitude, and measure perceptual fidelity, layout, mixed tilt, and prompt retention jointly.

Generative vocabulary. This is the decisive axis. We propose segmenting many fitted fields into persistent trajectory tubes using motion and semantic features, learning a dictionary whose entries aggregate across source videos, and holding out entire motions, appearances, and compositions. A speaker should emit prototype identities plus continuous transforms and residuals; a local flow or diffusion head can model residual 22-D marks [9, 10, 4]. The conclusive pitch experiment is not merely prettier reconstruction. It is a prompt-only render whose every macro token is source-independent, whose prompt and motion are held out compositionally, and whose correct/shuffled/null margins persist over many classes and seeds.

7 Limitations

Closed semantic domain. The prompt test contains three classes and 18 source fields. It does not measure open vocabulary, long-form prompting, camera control, or physical interaction.

Source-backed macro vocabulary. Foreground and background trajectories reuse programs derived from fitted sources. Although no whole video is retrieved and two donors are composed, the system has not learned source-independent motion concepts. This is the principal limitation.

Small, synthetic evaluation. OpenCLIP retrieval is noisy at the current visual quality and does not replace human judgment, FVD [17], or a broad action benchmark. The reported confidence intervals cover seeds within fixed small datasets, not dataset uncertainty.

Rendering and color. Additive overlap can wash out or overshoot RGB range, and the renderer does not model occlusion through ordered alpha compositing. Support truncation is exact only with respect to the declared five-sigma approximation.

No demonstrated efficiency advantage. We do not claim that a Jewel video is smaller, cheaper, or higher quality than a latent video. A 72k-by-22 continuous field is not intrinsically a compression format. The motivation is persistent executable structure; editing and scaling advantages remain hypotheses requiring direct baselines.

Negative experiments are local evidence. Failure of a particular global decoder, local phrase model, or learned retrieval gate identifies a missing dependency under its tested data and optimizer. A single failed run is not treated as a law about Gaussian languages.

8 Broader impact

An explicit spacetime program could enable inspectable animation tools, localized corrections, and reuse of persistent state. If future models make local edits without regenerating an entire video, they may reduce iteration cost. The same capability can enable deceptive synthetic media, impersonation, or biased visual depictions; explicit primitives do not remove those harms. Before releasing a capable open-ended generator, appropriate measures include provenance metadata, watermarking where robust, prompt and output abuse screening, consent-aware data curation, bias and identity audits, and staged access. Training and rendering also consume energy. Efficiency claims should be based on measured end-to-end quality-matched compute, not primitive counts alone.

9 Conclusion

The evidence supports a specific proposition: persistent time-tilted Gaussian Jewels form a useful video basis, their amortized prediction improves with data, their physical parameters admit an irregular hierarchical tokenization, and prompt plus seed can drive a bounded native program with controlled semantic selectivity. It does not yet support the claim that we have a general text-to-video model. The gap is sharply defined rather than mysterious: replace source-backed trajectory macros with reusable learned prototypes, then demonstrate compositional prompt-only generation on held-out motions and appearances. That experiment is the shortest path from a convincing feasibility result to a compute-and-data scaling proposal.

References

- [1] Andreas Blattmann, Tim Dockhorn, Sumith Kulal, Daniel Mendelevitch, Maciej Kilian, Dominik Lorenz, Yam Levi, Zion English, Vikram Voleti, Adam Letts, Varun Jampani, and Robin Rombach. Stable video diffusion: Scaling latent video diffusion models to large datasets. *arXiv preprint arXiv:2311.15127*, 2023. URL <https://arxiv.org/abs/2311.15127>.
- [2] Chameleon Team. Chameleon: Mixed-modal early-fusion foundation models. *arXiv preprint arXiv:2405.09818*, 2024. URL <https://arxiv.org/abs/2405.09818>.
- [3] Jonathan Ho, Tim Salimans, Alexey Gritsenko, William Chan, Mohammad Norouzi, and David J. Fleet. Video diffusion models. In *Advances in Neural Information Processing Systems*, volume 35, 2022. URL <https://arxiv.org/abs/2204.03458>.
- [4] Ka-Hei Hui, Chao Liu, Xiaohui Zeng, Chi-Wing Fu, and Arash Vahdat. Not-so-optimal transport flows for 3d point cloud generation. *arXiv preprint arXiv:2502.12456*, 2025. URL <https://arxiv.org/abs/2502.12456>.
- [5] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkuehler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics*, 42(4), 2023. doi: 10.1145/3592433.
- [6] Inseo Lee, Youngyoon Choi, and Joonseok Lee. Gaussianvideo: Efficient video representation and compression by gaussian splatting. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, 2025. URL https://openaccess.thecvf.com/content/CVPR2025W/PBVS/html/Lee_GaussianVideo_Efficient_Video_Representation_and_Compression_by_Gaussian_Splatting_CVPRW_2025_paper.html.

- [7] Zhan Li, Zhang Chen, Zhong Li, and Yi Xu. Spacetime gaussian feature splatting for real-time dynamic view synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. URL https://openaccess.thecvf.com/content/CVPR2024/html/Li_Spacetime_Gaussian_Feature_Splatting_for_Real-Time_Dynamic_View_Synthesis_CVPR_2024_paper.html.
- [8] Chenguo Lin, Panwang Pan, Bangbang Yang, Zeming Li, and Yadong Mu. Diffsplat: Repurposing image diffusion models for scalable gaussian splat generation. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2501.16764>.
- [9] Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.02747>.
- [10] David Lüdke, Enric Rabasseda Raventós, Marcel Kollovieh, and Stephan Günnemann. Unlocking point processes through point set diffusion. *arXiv preprint arXiv:2410.22493*, 2024. URL <https://arxiv.org/abs/2410.22493>.
- [11] Panwang Pan, Chenguo Lin, Chenxin Li, Jingjing Zhao, Yuchen Lin, Haopeng Li, Yunlong Lin, Kairun Wen, Yixuan Yuan, and Yadong Mu. Diff4splat: Repurposing video diffusion models for dynamic scene generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4232–4244, 2026.
- [12] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139, pages 8748–8763, 2021.
- [13] Barbara Roessle, Norman Mueller, Lorenzo Porzi, Samuel Rota Buló, Peter Kotschieder, Angela Dai, and Matthias Niessner. L3dg: Latent 3d gaussian diffusion. *ACM Transactions on Graphics*, 2024. URL <https://arxiv.org/abs/2410.13530>.
- [14] Sachin Shah, Anustup Choudhury, Guan-Ming Su, Jaclyn Pytlarz, Christopher A. Metzler, and Trisha Mittal. Gaussian representations for video. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 827–837, 2026. URL https://openaccess.thecvf.com/content/WACV2026/html/Shah_Gaussian_Representations_for_Video_WACV_2026_paper.html.
- [15] Weronika Smolak-Dyżewska, Dawid Malarz, Kornel Howil, Jan Kaczmarczyk, Marcin Mazur, and Przemysław Spurek. Vegas: Video gaussian splatting. *arXiv preprint arXiv:2411.11024*, 2024. URL <https://arxiv.org/abs/2411.11024>.
- [16] Yang-Tian Sun, Yi-Hua Huang, Lin Ma, Xiaoyang Lyu, Yan-Pei Cao, and Xiaojuan Qi. Splatter a video: Video gaussian representation for versatile processing. In *Advances in Neural Information Processing Systems*, volume 37, pages 50401–50425, 2024.
- [17] Thomas Unterthiner, Sjoerd van Steenkiste, Karol Kurach, Raphael Marinier, Marcin Michalski, and Sylvain Gelly. Towards accurate generative models of video: A new metric & challenges. *arXiv preprint arXiv:1812.01717*, 2018.
- [18] Longan Wang, Yuang Shi, and Wei Tsang Ooi. Gsvc: Efficient video representation and compression through 2d gaussian splatting. *arXiv preprint arXiv:2501.12060*, 2025. URL <https://arxiv.org/abs/2501.12060>.
- [19] Xinlong Wang, Xiaosong Zhang, Zhengxiong Luo, Quan Sun, Yufeng Cui, Jinsheng Wang, Fan Zhang, et al. Emu3: Next-token prediction is all you need. *arXiv preprint arXiv:2409.18869*, 2024. URL <https://arxiv.org/abs/2409.18869>.
- [20] Guanjun Wu, Taoran Yi, Jiemin Fang, Lingxi Xie, Xiaopeng Zhang, Wei Wei, Wenyu Liu, Qi Tian, and Xinggang Wang. 4d gaussian splatting for real-time dynamic scene rendering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 20310–20320, 2024.

- [21] Bowen Zhang, Jiapeng Tang, Matthias Niessner, and Peter Wonka. Gaussiancube: A structured and explicit radiance representation for 3d generative modeling. *arXiv preprint arXiv:2403.19655*, 2024. URL <https://arxiv.org/abs/2403.19655>.
- [22] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 586–595, 2018.
- [23] Chunting Zhou, Lili Yu, Arun Babu, Kushal Tirumala, Michihiro Yasunaga, Leonid Shamis, Jacob Kahn, Xuezhe Ma, Luke Zettlemoyer, and Omer Levy. Transfusion: Predict the next token and diffuse images with one multi-modal model. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2408.11039>.

A Protocols and preregistered gates

Temporal-tilt gate. The inventory contains three source videos and seeds 0, 1, and 2. Free and projected models share primitive and serialized parameter counts. The projection zeroes covariance entries (u, t) and (v, t) after each optimizer step. The gate requires a positive mean paired PSNR difference, a 95% interval excluding zero, and removal of measured mixed tilt in the control.

Encoder-scaling gate. Nested train sets contain 12, 60, and 120 examples. The common validation set contains 60 examples and has SHA-256 digest 3e034a1263e21b566c79b63a8257e9794c4c66316e3a82d6fe65d418d1699b25. The encoder uses 4,096 sampled query points per step, AdamW, an initial matched 120-corpus-pass phase, and a continuation at learning rate 3×10^{-5} for at least 40 and at most 80 additional passes. Evaluation occurs every 10 passes, with patience 3 and minimum improvement 0.03 dB. The reported endpoint is the common continuation protocol, not per-arm cherry-picked training length.

Physical-language gate. The active books contain 1024 entries for covariance, 1024 for base/surface color, and 1024 for the color Jacobian. Centers are not quantized. Required checks include finite metrics, token-only PSNR above 20 dB, mixed-tilt retention above 0.8, exactly zero grid locking, fewer than 40k decisions for eight frames, and same-source canonicity margin above 0.05.

Prompt-only gate. The exact prompts are a ballerina spinning a pirouette in a studio, a golden retriever catching a ball on grass, and a welder joining steel with bright sparks. Seeds are 20260914, 20260915, and 20260916. Controls cyclically rotate prompts or remove text while holding the declared seed. The learned speaker trains on each exact prompt and two paraphrases; a fourth paraphrase per class is held out. Its AdamW learning rate is 0.002, batch size 64, text-dropout probability 0.1, minimum 1000 steps, maximum 5000 steps, evaluation interval 100, and plateau patience 10 evaluations with minimum NLL improvement 10^{-4} .

B Additional quantitative detail

Table 2: Encoder scaling audit. Intervals are 95% t intervals over three seed means.

Train videos	PSNR (dB)	95% CI	LPIPS	Layout PSNR
12	21.876	[21.751, 22.000]	0.4728	29.812
60	22.399	[22.344, 22.453]	0.4651	30.726
120	22.628	[22.459, 22.797]	0.3916	31.533
Teacher	27.699	—	0.1720	—

The learned speaker’s higher absolute similarity is not evidence that it passes the stricter gate: the relevant outcome is relative retrieval among candidate prompts for each generated sample. It passes correct-over-shuffled in 7/9 cases but strict top-1 in only 4/9.

Table 3: Prompt-only semantic audit. Similarities are OpenCLIP cosine values.

Speaker	Correct	Shuffled	Null
Exact compiler	0.18173	0.12525	0.14825
Learned, unseen paraphrases	0.25056	0.20274	0.20939

C Compute and implementation

Experiments ran on local NVIDIA RTX 4090 (24 GB) and RTX 2070 Super (8 GB) GPUs. Support-complete 72k-field rendering measured 8.01 GB peak allocation on the RTX 4090; chunked paths permit 45k–120k field fitting on the 2070 Super. The encoder has 5.26M trainable parameters. Representative 12-example encoder continuations take approximately 129 seconds for 2,400 updates on the recorded machine; larger arms scale with corpus size and evaluation cost. The learned speaker has 541,223 parameters and trains for 1,100 updates in the selected run. Derivative appearance studies run to 12k–16k updates. Total exploratory compute exceeds these selected runs because architectural screening, renderer validation, and failed controls are retained in the repository rather than omitted. A submission artifact should record exact GPU model, driver, package lock, wall time, and energy estimate for every reproduced run.

D Reproducibility map

The following repository-relative machine-readable reports support the manuscript:

- `sol/results/temporal_tilt_replication_v1/report.json`;
- `sol/results/encoder_convergence_v2_continued/aggregate/report.json` and `audit/audit.json`;
- `sol/results/jewel_casting_language_v0/hierarchical_v1/gate0f_individual/report.json`;
- `sol/results/jewel_casting_language_v0/prompt_trajectory_speaker_v1/report.json`;
- `sol/results/jewel_casting_language_v0/learned_trajectory_speaker_v1/report.json` and `audit/report.json`; and
- `sol/results/local_adapter_convergence_v1/README.md`.

Reports include protocol fields, artifact ownership, per-sample metrics, checks, and declared gate outcomes. Code and reports are retained, but an anonymized public artifact with a complete license audit has not yet been prepared; this is recorded as No in the checklist rather than implied.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: Sections 1 and 7 state that this is a bounded concept-and-feasibility result, distinguish exact and learned speakers, and identify the source-backed macro dictionary as the central unsolved limitation.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.

- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 discusses vocabulary size, source ownership, evaluator scope, appearance quality, absent second-window carry, data scale, and synthetic-media risks.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper introduces model definitions and empirical hypotheses but no theorem or proof claim.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Appendices A and D specify splits, gates, retained reports, seeds, counts, renderer settings, and commands for the principal experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: All code and machine-readable reports are retained, but an anonymized public artifact bundle has not yet been prepared; this must be completed before submission.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 and Appendix A give the data ownership, model size, training conditions, controls, and evaluation rules needed to interpret the results.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Table 1 reports three-seed intervals for encoder scaling and a paired 95% interval over nine matched temporal-tilt runs; exact prompt outcomes report all nine programs rather than only averages.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix C records the RTX 4090 target, the RTX 2070 compatibility path, parameter counts, fitting time, and representative training budgets. Total exploratory compute is acknowledged as larger than the selected runs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses synthetic or established benchmark video, reports scope and provenance, and does not involve human subjects or sensitive personal data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 8 discusses editable visual tools and efficiency as potential benefits, and deception, impersonation, bias, and compute use as potential harms.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No open-ended high-capability generator or scraped dataset is released with this feasibility draft. Section 8 describes safeguards required before a future model release.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: The paper credits the existing models and datasets used, but a submission-ready license table and redistribution audit for every retained synthetic derivative are still pending.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not introduce a public dataset or externally released pretrained model; retained experimental artifacts are reproducibility records.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The work contains no crowdsourcing experiment and no research with human participants.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The work contains no human-subject research requiring institutional review.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

711 Answer: [N/A]
712 Justification: No pretrained LLM or agent is an original component of the scientific method.
713 Automated assistance was used for ordinary code and manuscript preparation, which the
714 policy does not require declaring.
715 Guidelines:
716 • The answer [N/A] means that the core method development in this research does not
717 involve LLMs as any important, original, or non-standard components.
718 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
719 be described.